

PHYS-512: Problem Set 4

Simon Briesenick

October 15, 2022

- (a) Newton's method is implemented similarly to the in-class example. The resulting fit for 50 iterations is shown in Fig. 1a. To implement Newton's method, one needs the gradient of the fit function




Figure 1: Single Lorentzian Fit to the signal and associated error in the number of fitting iterations.

w.r.t. the model parameters m :

```
def lorentzian(t, m):
    a, t_0, w = m
    d = a/(1 + ((t - t_0)**2)/(w**2))
    # make empty array to store gradient values
    grad = np.zeros([t.size, 3])
    # derivative w.r.t. a
    grad[:,0] = (w**2)/((w**2) + (t - t_0)**2)
    # deriv t_0
    grad[:,1] = 2*a*w**2*(t - t_0)/(((w**2) + (t - t_0)**2)**2)
    # deriv w
    grad[:,2] = 2*a*w*(t - t_0)**2/(((w**2) + (t - t_0)**2)**2)
    return d, grad
```

A for loop then solves the algebraic system

$$A_m^T N^{-1} A_m \delta m = A_m^T N^{-1} r \quad (1)$$

for δm - which serves as a correction to each guess, i.e.

$$m \rightarrow m + \delta m. \quad (2)$$

The implementation is very similar to how it was done in class:

```
model_params[0] = (a, t_0, w)
error = []
for i in range(iterator):
    pred, grad = lorentzian(time, model_params[i])
    # calculate residual (serves as noise covariance matrix)
    r = d - pred
    err = np.mean(np.abs(r))
    error.append(err**2)
    # cast matrices in readily available forms
    r = np.matrix(r).transpose()
    grad = np.matrix(grad)
    # compute linear system components
    lhs = grad.transpose()*grad
    rhs = grad.transpose()*r
    # solve it for dm
    dm = np.linalg.inv(lhs)*rhs
    # output is matrix, reshape for next step
    dm = dm.reshape(3,)
    if i == iterator - 1:
        break
    else:
        model_params[i + 1] = model_params[i] + dm
```

As initial guess for the model parameters, values could easily be read of from a preliminary plot of the signal as (units dropped)

$$t_0 = 1.8 \times 10^{-4}, a = 1.5, w = 1 \times 10^{-4}.$$

The final values for the parameters are

$$t_0 = 1.92 \times 10^{-4}, a = 1.42, w = 1.79 \times 10^{-5}.$$

(b) To estimate the uncertainties on the calculated parameters, one needs to compute

$$\left(A^T N^{-1} A\right)^{-1} \quad (3)$$

and take the square root of the diagonal entries. The Code is very similar to how I have done it in the previous P-set:

```
par_errs = np.sqrt(np.diag(np.linalg.inv(lhs)*err**2))
```

For the parameter errors, I get

$$\delta t_0 = 4.1 \times 10^{-9}, \delta a = 3.3 \times 10^{-4}, \delta w = 5.8 \times 10^{-9}.$$

- (c) This exercise follows from (d), upon forcing the model parameters for the two additional Lorentzian functions to vanish (Originally, I just took a `single_lorentzian(t, m)` for modeling the data. The Code is only marginally different and is not repeated here). The parameters are the approximately as before:

$$t_0 = 1.92 \times 10^{-4}, a = 1.42, w = 1.80 \times 10^{-5}$$

to two significant figures.

The absolute difference in a parameter p , $\Delta p = |p_{\text{analytic}} - p_{\text{numeric}}|$ with the two methods is calculated as

$$\Delta t_0 = 1.6 \times 10^{-8}, \Delta a = 3.0 \times 10^{-3}, \Delta w = 2.9 \times 10^{-8}.$$

However, it should be noted that the discrepancies in those differences depend heavily on the chosen step size. The above results are obtained with `step = 1e-6` (as used in (d)), while `step = 1e-10` leads to much better agreement

$$\Delta t_0 = 1.0 \times 10^{-12}, \Delta a = 1.1 \times 10^{-7}, \Delta w = 2.9 \times 10^{-12}.$$

But in either case, the values lie within each others bounds and the deviations are not statistically significant. It was shown in P-set 1 (or rather in class leading up to PS1), that an optimal step-size for the forward difference method is roughly proportional to $\sqrt{\epsilon_f}$ under the assumption that f and f'' are approximately equal to each other. Under this assumption an ideal step-size is calculated to be on the order of `step = 1e-8`. Clearly, there is no single `step` to minimize all truncation errors at the same time, since each second-order partial derivative behaves differently. In principle, one could minimize the truncation errors further by using separate `steps` for the individual parameters. It seems that the relative differences are therefore statistically significant, only when a sub-ideal choice for `step` is used.

- (d) At the heart of this implementation lies the forward difference method for each model parameter m_i :

```
[TRUNCATED]
# numerical derivs
step = 1e-6
for i, v in enumerate(m):
    dm = np.zeros(len(m))
    dm[i] = step
    grad[:,i] = (f(m + dm) - f(m))/step
```

where `grad` has the same functionality as above. `f(m)` represents the given model that calls three instances of `single_lorentzian(t, m)`

```
def single_lorentzian(t, m):
    return m[0]/(1 + ((t - m[1])**2)/(m[2]**2))
def grad_lorentz(t, m):
    [TRUNCATED]
    fun = single_lorentzian
    # define individual instance of Lorentzian
    f_single = lambda amp, center, width : (
        fun(t, np.array([amp, center, width]))
```

```

    )
    # triple Lorentzian model
    f = lambda m : (
        f_single(m[0], m[3], m[5]) +
        f_single(m[1], m[3] + m[4], m[5]) +
        f_single(m[2], m[3] - m[4], m[5]))

```

lambda Notation is useful here, since calling the functions is easier and the code is overall more readable. Using this method, the model parameters are found to be

$$\begin{array}{lll}
 a = 1.44 & b = 6.53 \times 10^{-2} & c = 1.05 \times 10^{-1} \\
 t_0 = 1.93 \times 10^{-4} & dt = 4.44 \times 10^{-5} & w = 1.60 \times 10^{-5}
 \end{array}$$

with corresponding uncertainties (calculated similarly as described in (b))

$$\begin{array}{lll}
 \delta a = 2.3 \times 10^{-4} & \delta b = 2.2 \times 10^{-4} & \delta c = 2.2 \times 10^{-4} \\
 \delta t_0 = 2.7 \times 10^{-9} & \delta dt = 3.2 \times 10^{-8} & \delta w = 4.9 \times 10^{-9}
 \end{array}$$

Complete Code:

```

def single_lorentzian(t, m):
    return m[0]/(1 + ((t - m[1])**2)/(m[2]**2))

def grad_lorentz(t, m):
    # model params m
    # m[0] = a      m[3] = t_0
    # m[1] = b      m[4] = dt
    # m[2] = c      m[5] = w
    fun = single_lorentzian
    f_single = lambda amp, center, width : (
        fun(t, np.array([amp, center, width]))
    )
    f = lambda m : (
        f_single(m[0], m[3], m[5]) +
        f_single(m[1], m[3] + m[4], m[5]) +
        f_single(m[2], m[3] - m[4], m[5]))
    pred = f(m)
    # make empty array to store gradient values
    grad = np.empty([t.size, 6])
    # numerical derivs
    step = 1e-6
    for i, v in enumerate(m):
        dm = np.zeros(len(m))
        dm[i] = step
        grad[:,i] = (f(m + dm) - f(m))/step
    return pred, grad

```

Fig. 2 shows fitting the data to a sum of three Lorentzian functions. This model compares better with those used in (a)/(c). One can compare the RSS values after some number of fit iterations to see that the initial implementation led to $RSS \approx 60$ while this one leads to $RSS \approx 20$.

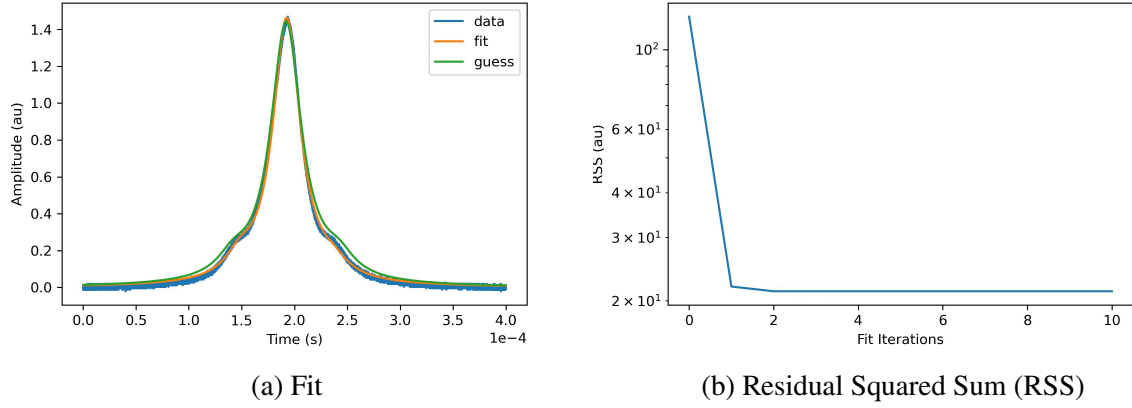


Figure 2: Single Lorentzian Fit to the signal and associated error in the number of fitting iterations.

- (e) If the noise was uncorrelated, we would expect the residual to randomly scatter around the center line. As seen in Fig. 3, this is clearly not the case.

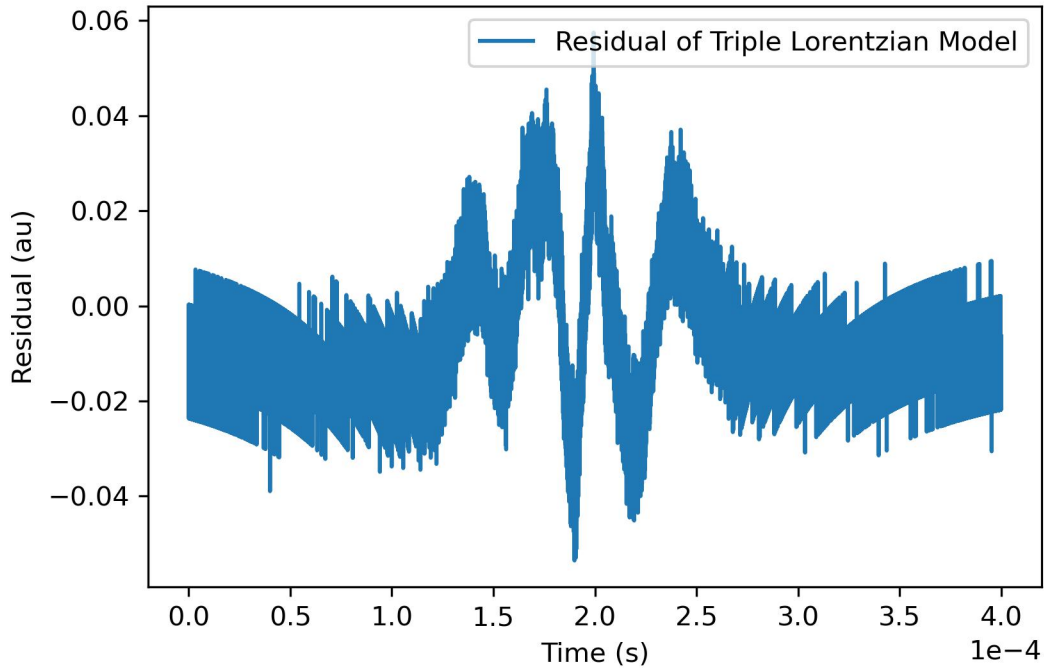


Figure 3: Residual in triple Lorentz model from (d).

- (f) The following code generates realizations of models with parameters offset from their best-fit values as found in (d) from the full covariance matrix.

```
# calculate covariance matrix
cov = np.linalg.inv(lhs)*err
sampler = 6
```

```

for j in range(sampler):
    alt_params = np.random.multivariate_normal(model_params[-1], cov)
    plt.plot(time, grad_lorentz(time, alt_params)[0],
             label=f"perturbation {j + 1}")

```

The offset is determined from a 6D normal distribution of the parameters around their best fit values. The diagonal entries are the variances of the parameters "on their own", while the off-diagonals are the co-variances and determine how two parameters are correlated, i.e. if one were to take a slice of this multi-dimensional normal distribution parallel to two parameter axis, the amount by which a sphere (in the uncorrelated case) is distorted into an ellipse is given by the off-diagonal entries. The results are shown in Fig. 4. The best-fit, as well as all perturbations fail to "correctly" reproduce




Figure 4: Perturbing the best fit parameters by offsets set through the covariance matrix gives separate implementations. 4b shows a magnified cut-out near the central maximum.

the data. This is especially apparent around the central maximum. The relative χ^2 values, $\Delta\chi^2$, are shown in Fig. 5. On average, $\Delta\chi^2 \approx 0.08$.

- (g) To run an MCMC chain for this fitting problem, I first imported the results from the previous exercises - most importantly the best-fit parameters, as well as the associated covariances:

```

import p_set4d_e as prev
import numpy as np
from matplotlib import pyplot as plt

# from previous answer, generate step size array and initial values
par_best= prev.model_params[-1]
par_errs = np.sqrt(np.diag(np.linalg.inv(prev.lhs)))

```

Since it will cancel out in later computation anyway, I dropped the calculation of σ^2 . To evaluate the quality of a specific run in the MCMC chain, a helper function for the computation of χ^2 is implemented

```

def get_chi_sq(data, params):
    pred, grad = grad_lorentz(time, params)

```




Figure 5: $\Delta\chi^2$ for 10,000 realizations.

```
chisq = np.sum((data - pred)**2)
return chisq
```

The initial parameters are taken from the best-fit values and their covariances:

```
n_steps = 100000
steps_taken = 0
chi_sq = []
par = []
# initialize parameter list by using best fit +/- a little something
init_par = par_best + 3.0*np.random.randn(len(par_best))*par_errs
chi_sq.append(get_chi_sq(d, init_par))
par.append(init_par)
```

The MCMC chain is then started for `n_steps` runs. The result is shown below. The typical run-in is followed by a fast convergence to an optimum. To visualize the parameter distribution in the chain, corner plots are shown in Fig. 7 with corresponding central values for each of the parameters, as well as their errors. While the error bars stayed approximately the same, the parameter values did change - even significantly for b which jumped from $b_{\text{Newton}} = 6.47 \times 10^{-2}$ to $b_{\text{MCMC}} = 6.27 \times 10^{-2}$. Complete Code:

```
import p_set4d_e as prev
import numpy as np
from matplotlib import pyplot as plt
# from previous answer, generate step size array and initial values
par_best = prev.model_params[-1]
par_errs = np.sqrt(np.diag(np.linalg.inv(prev.lhs)))
```




Figure 6: χ^2 (unscaled) for MCMC with 100,000 cycles.

```
def get_chi_sq(data, params):
    pred, grad = grad_lorentz(time, params)
    chisq = np.sum((data - pred)**2)
    return chisq

n_steps = 100
steps_taken = 0
chi_sq = []
par = []
# initialize parameter list by using best fit +/- a little something
init_par = par_best + 3.0*np.random.randn(len(par_best))*par_errs
chi_sq.append(get_chi_sq(d, init_par))
par.append(init_par)

while steps_taken < n_steps:
    # generate new parameters
    par_new = par[-1] + 1.0*np.random.randn(len(par_best))*par_errs
    # calculate new chi squared
    new_chi = get_chi_sq(d, par_new)
    delta_chi = new_chi - chi_sq[-1]
    prob_step = np.exp(-0.5*delta_chi)
```



```

# if prob_step is greater than one, chi will have increased
# and we want to discard that step, otherwise accept with a certain
# probability:
if prob_step > np.random.rand(1):
    par.append(par_new)
    chi_sq.append(new_chi)
else:
    # if step is not taken, copy last entries for params and chi_sq
    par.append(par[-1])
    chi_sq.append(chi_sq[-1])
steps_taken += 1

```

(h) I calculate the cavity width with the equation

$$x_{\text{cavity}} = \frac{dt}{w}(9\text{GHz}) \quad (4)$$

to be $x_{\text{cavity}} = 250\text{GHz}$ with a fractional uncertainty of 0.001, i.e. an absolute uncertainty of 0.25GHz. I haven't actually derived this equation, but taken it from a class colleague with my own parameters.

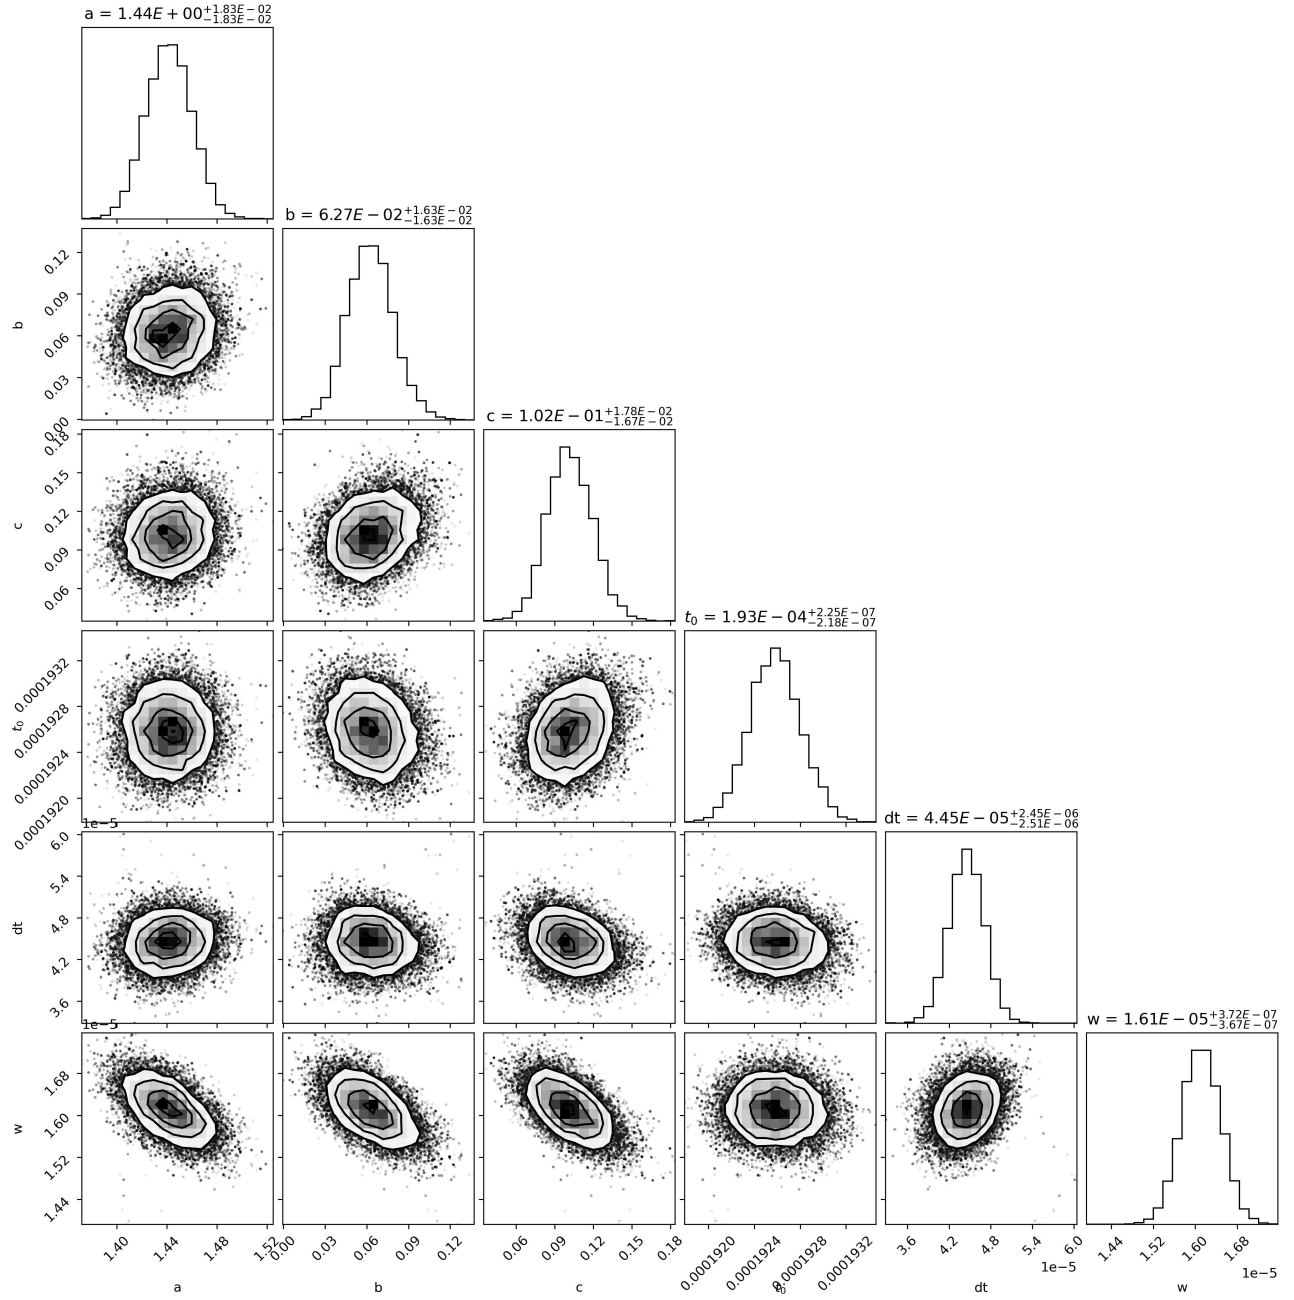


Figure 7: Corner Plots